

Priya Raghunathan

Location: Toronto, Canada

Email: priya.raghunathan@example.com

Phone: +1 416 555 0142

GitHub: github.com/example

OBJECTIVE

Data Scientist with 6 years turning messy operational data into decisions people actually act on. Strongest where modelling meets deployment: I ship models into production and stay responsible for them afterwards. Looking for a team that measures data science by outcomes rather than notebooks.

PROJECT HIGHLIGHTS

Churn Model That Changed the Retention Playbook

Retention spend was spread evenly across all at-risk accounts because nobody could rank them. Built a gradient-boosted churn model on **18 months** of billing and support data, calibrated so the score could be read as a probability rather than a rank. Shipped it behind an API the CRM calls nightly. Targeted campaigns on the top decile cut churn in that group by **19%** against a held-out control, on **~40k monthly accounts**.

Forecast Backtesting Harness

Demand forecasts were evaluated by eyeballing charts, so nobody could say whether a change helped. Built a backtesting harness that replays any candidate model over historical windows and reports error by horizon and by segment. It caught a seasonal regression that had been live for **two quarters**, and made model changes an argument about numbers rather than opinion.

WORK EXPERIENCE

Meridian Retail Analytics

Data Scientist | Apr 2021 - Present

- Own the churn and demand-forecasting models end to end: **Python, scikit-learn, XGBoost, Airflow**.
- Serve models through a **FastAPI** service on **Kubernetes**, with drift monitoring and weekly retraining.
- Work directly with the retention team to design experiments and read the results honestly.
- Cut the forecasting pipeline runtime from **6 hours to 40 minutes** by rewriting joins in **DuckDB**.

Northline Logistics

Data Analyst, then Data Scientist | Jul 2018 - Mar 2021

- Built the reporting layer in **dbt** and **BigQuery** that replaced a spreadsheet-based weekly pack.
- Developed a route-anomaly detector that flagged **~200 suspect trips/month** for review.
- Ran the A/B analysis for pricing changes across **12 regional markets**.

SKILLS

Modelling

- **Languages:** Python, SQL, R.
- **Libraries:** scikit-learn, XGBoost, statsmodels, pandas, Polars.
- **Methods:** Causal inference and A/B testing, time-series forecasting, calibration, survival analysis.

Engineering

- **Pipelines:** Airflow, dbt, BigQuery, DuckDB, Spark.
- **Deployment:** FastAPI, Docker, Kubernetes, MLflow, drift monitoring.
- **Visualisation:** Plotly, Streamlit, Looker.

EDUCATION

University of Waterloo

MMath Statistics | 2016 - 2018

University of Delhi

BSc Mathematics | 2013 - 2016